

AI unit 1

PAGE NO.

DATE

Turing Test and Rational Agent Approaches

state space Representation of Problems

Heuristic Search Techniques

Game Playing \rightarrow Min Max Search

Alpha Beta Cut off Procedures

AI is the branch of computer science that deals with creating machines or software that can think, learn, and act intelligently like humans.

AI \rightarrow giving machines the ability to perform tasks that normally requires human intelligence.

Key features of AI

Perception : Understanding Input from environment (e.g. vision, speech)

Reasoning : Drawing conclusions and making decisions

Learning : Improving performance from past experience (ML)

Problem-solving : Finding solutions using search, heuristics, or planning

Acting : Taking rational actions in real-world environments

Types of AI

\rightarrow based on capabilities

(i) Narrow AI (weak AI) : specialized for a single task or specific task or a narrow range of tasks.

- Highly specialized and operate within a limited context.
- Cannot generalize their knowledge or apply it to tasks outside their designated function

Ex. Chatbots, Google Maps, Face recognition

(ii) General AI (strong AI) : can perform any intellectual task that a human can do with the ability to generalize knowledge and apply it to different contexts.

AI is the science of making machines act intelligently
— to perceive, reason, learn, and decide like humans.

PAGE NO.	
DATE	

- It possesses the ability to understand, learn and apply knowledge across a wide range of tasks.
 - Broad intelligence → it can perform a variety of tasks, making it versatile and adaptable.
 - Human-like reasoning
 - Self-learning → capable of learning and improving over time, adapting to new situations and ~~and~~ acquiring new skills without human intervention
- Ex. General AI is theoretical, not yet achieved fully, but research is ongoing.

(iii) Super AI : Most advanced form of AI, surpassing human intelligence

• Super AI is still a theoretical concept

Examples of AI Applications

- Virtual assistants (Siri, Alexa)
- Self driving cars
- Medical diagnosis
- Game playing (Chess, Go)
- Fraud detection in Banking.

Test is when a person can't tell whether ans is given by machine or human

Turing Test

- Proposed by Alan Turing (1950)
- Turing asked "Can machines think?"

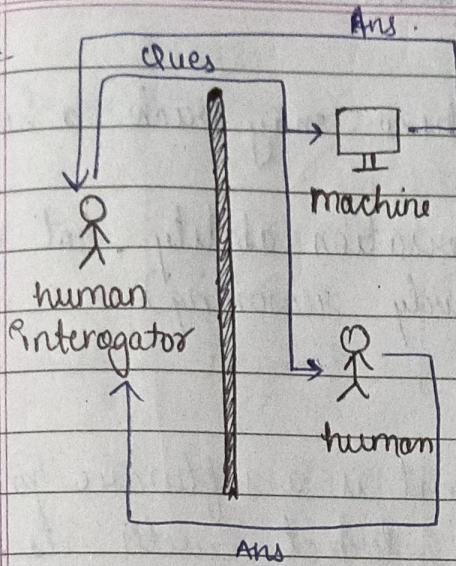
Instead of directly answering, he proposed an "imitation game" → known as Turing Test

Turing test is a widely recognized benchmark for evaluating a machine's ability to demonstrate human-like intelligence

Intelligence \rightarrow ability to exhibit human-like behavior.

PAGE NO.

DATE



Setup: A human interrogator communicates via text with two participants

1. A human

2. A machine (computer)

Both try to convince the interrogator that they are human.

Goal: If the interrogator cannot reliably tell the machine from the human, the machine is said to have passed the Turing Test

Criteria for Turing Test

The test does not require the machine to be correct or logical in its response but rather to be convincing in simulating human conversation.

The test is fundamentally about deception — the machine must fool the judge into believing that it is human.

Key criteria:

(i) Natural Language Processing (NLP): machine must understand and generate human language fluently.

(ii) Knowledge Representation: Machine needs to handle and manipulate knowledge to provide contextually relevant responses.

(iii) Reasoning: Machine should demonstrate some form of logical reasoning, even if flawed, to sustain a convo.

(iv) Learning: Machine should learn from its environment

Ex Judge: Are you a computer? \rightarrow No.

Multiply 29878978 to 199810498 \rightarrow (After long pause) wrong answer

Add 9187 with 6843 \rightarrow (After 20 sec) gives right answer

Turing Test is still philosophically important but not a practical benchmark today.

Modern AI is evaluated by performance metrics (accuracy, efficiency) not just by imitation

PAGE NO.	20
DATE	

Limitations

1. Deception over intelligence : A machine may trick a human without truly 'understanding'
2. Narrow measure : only tests conversation ability, not other forms of intelligence (vision, creativity, reasoning)

8.

Rational Agents in AI

An AI agent is a software program that can interact with its environment, collect data and use all data to perform self-directed tasks.

An agent is anything that perceives its environment through sensors and acts upon that environment using actuators

$$\boxed{\text{Agent} = \text{perception} + \text{Action}}$$

Key features of Agent

- Autonomous: Act without constant human input and decide next steps from past data
- Goal driven: Optimize for defined objectives
- Perceptive: Gather info from sensors, inputs or APIs
- Adaptable: Adjust strategies when situation change
- Collaborative: work with humans or other agents towards shared goals

Agent function : Maps percepts (inputs) to actions

$$f: P^* \rightarrow A$$

Agent program : Implementation of the agent function that runs on the agent's hardware

A rational agent is one that chooses actions that maximize its expected performance measure given the knowledge and percepts it has.

PAGE NO.	
DATE	

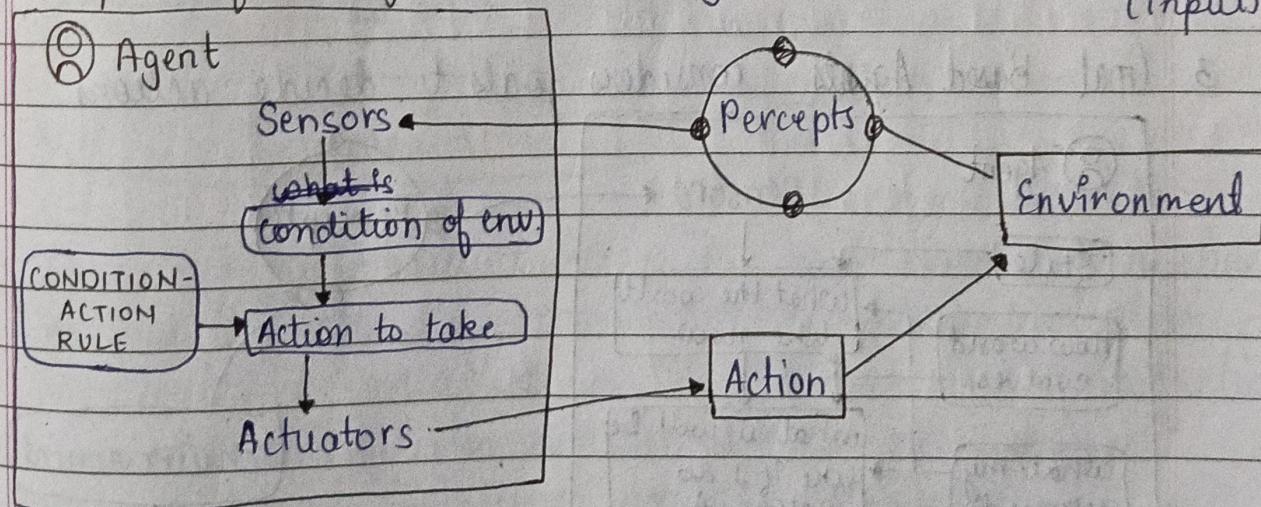
11

AI Agents Classification

- Reactive agents : respond to immediate environmental stimuli without foresight or planning
- Proactive agents : anticipate future states and plan actions to achieve long-term goals
- Single - agent system : one agent solves a problem independently
- Multi - agent systems : multiple agents interact, coordinate or compete to achieve goals ; may be homogeneous (same goals/roles) or heterogeneous (diverse roles)
- Rational agents : choose actions to maximize expected outcomes using both current and historical information.

Types of Agents

1. Simple Reflex Agents : Acts only on the current perception (input)



These agents respond directly to stimuli without considering past experiences or potential future states.

Uses condition-action rules ("If-then" rules)

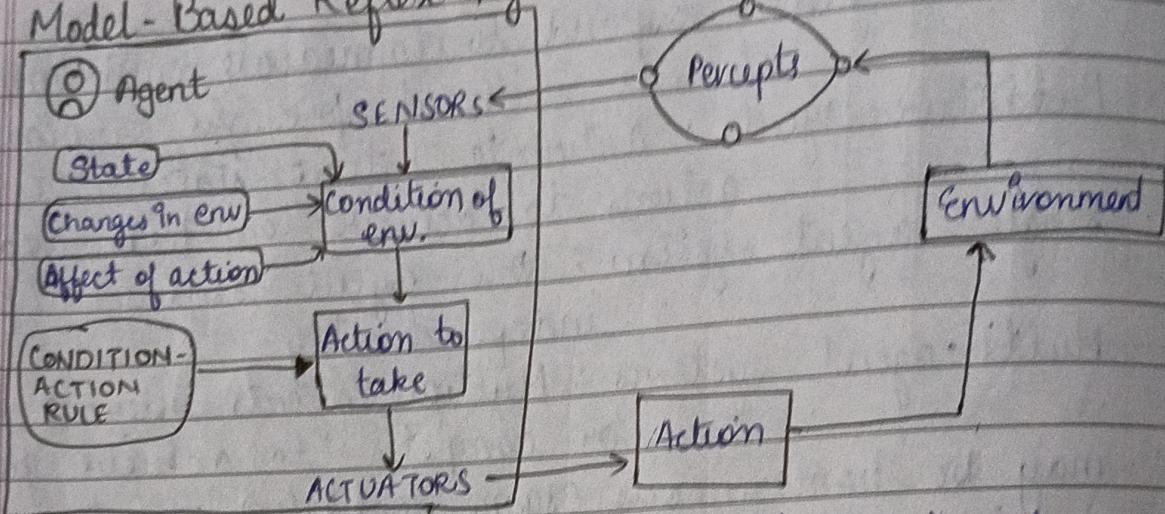
No memory of past states

Not very intelligent → works on fixed rules.

Ex - Traffic light controls on current traffic situation ☺

- Thermostat ("If temp < 18°C → then turn on heater")

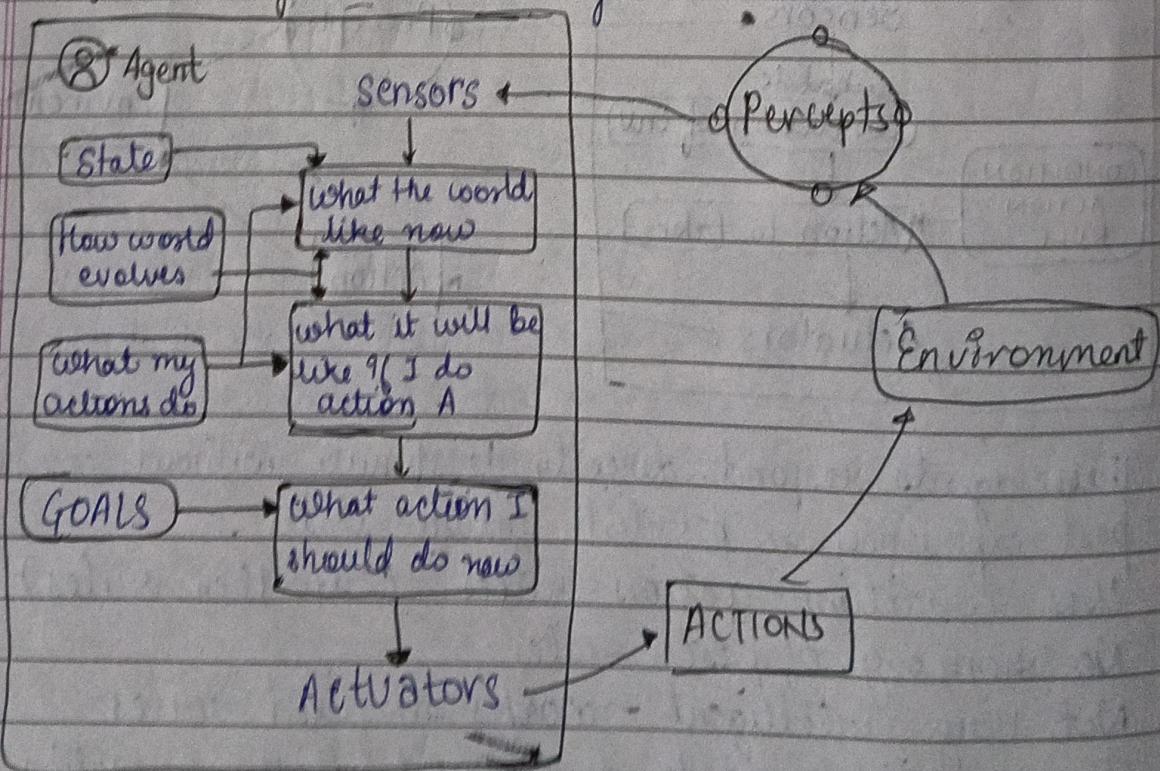
2 Model-Based Reflex Agent



- Maintains an internal model of world (env.)
 - function effectively in partially observable environments
- This internal model helps them make more informed decisions by considering how the world evolves and how their actions affect it.

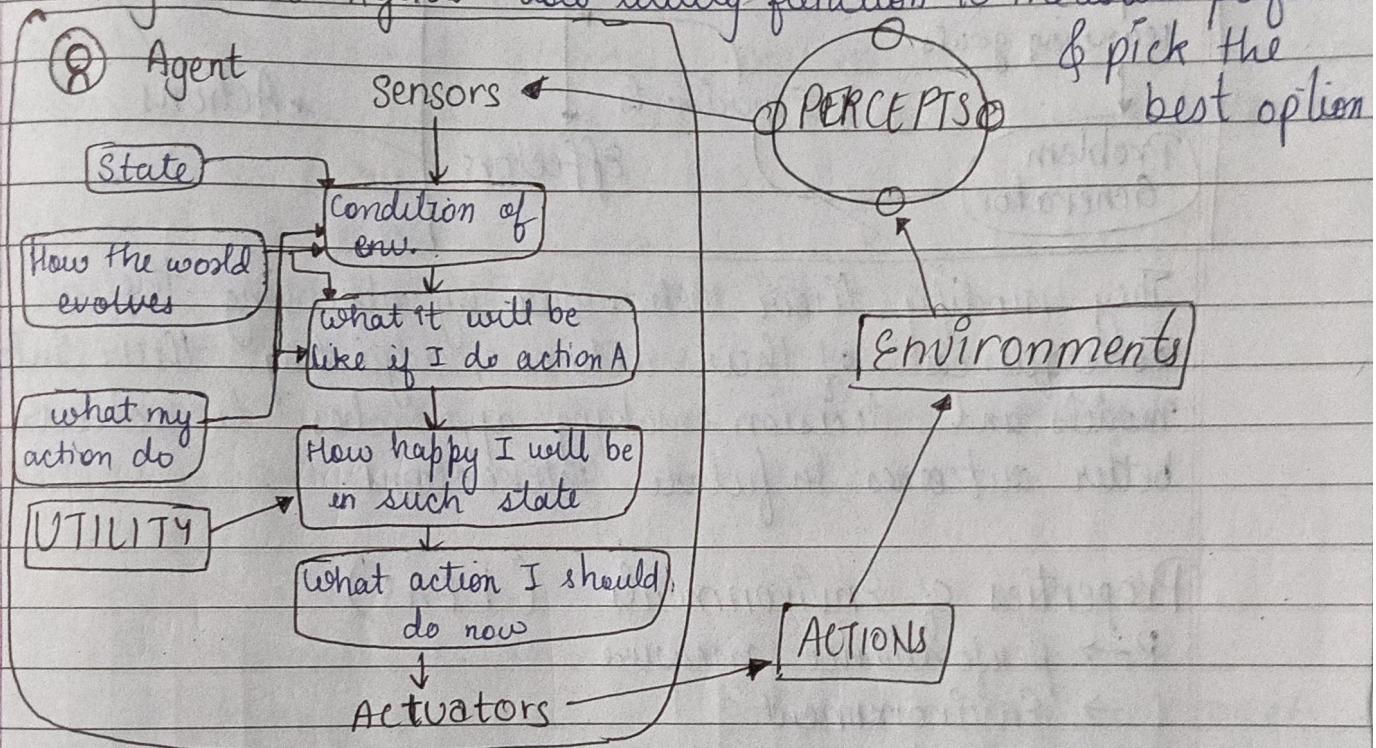
Ex. Robot vacuum cleaners that map rooms and tracks cleaned areas

3 Goal-Based Agents : considers goals to decide actions



- Not flexible → have fixed goal
- Evaluate how different action sequences might lead towards their defined goal, selecting the path that appears most promising
- Employ search and planning mechanisms
- May explore multiple possible routes to goal
Ex GPS navigation system choosing a path to destination

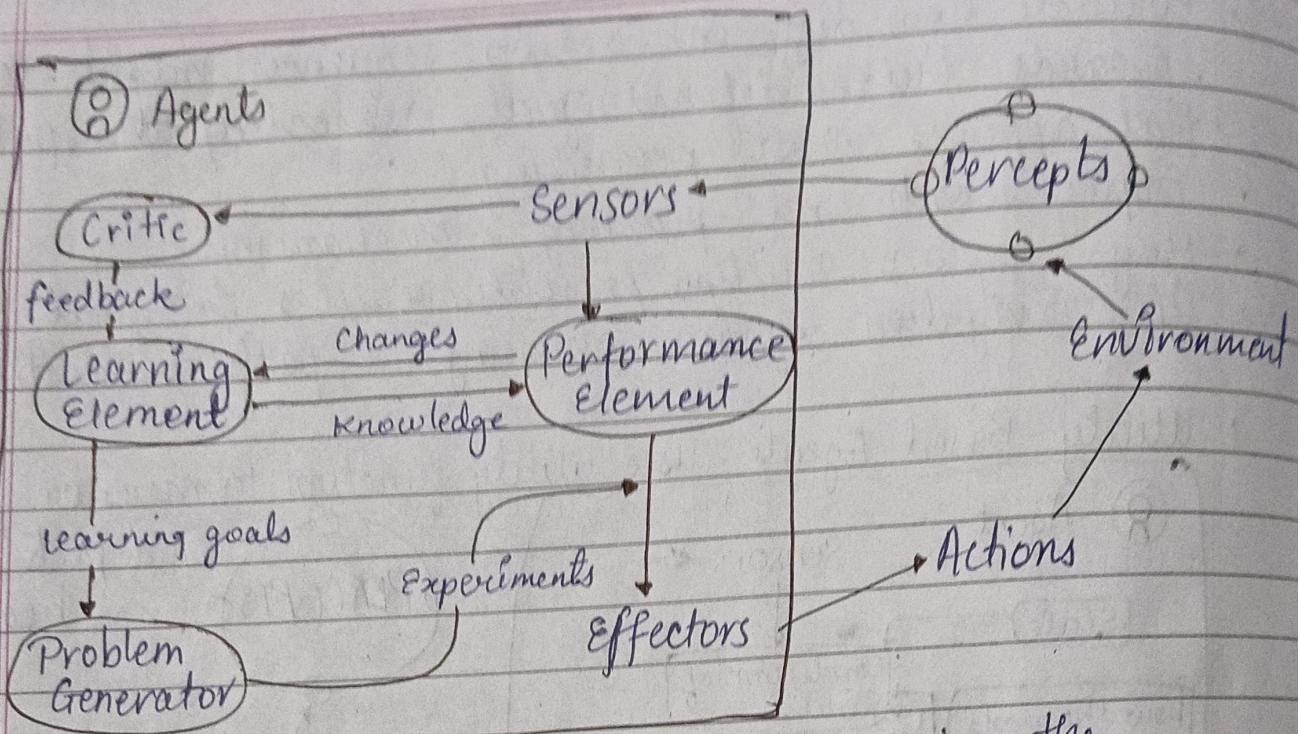
4. Utility-Based Agents : uses utility function to measure performance & pick the best option



- Extend goal-based thinking by evaluating actions based on how well they maximize a utility function — essentially a measure of "happiness" or "satisfaction".
- Handle probabilistic and uncertain environments
- Make rational decisions under constraints
- Ex Investment AI maximizing profit, minimizing risk.

5. Learning Agents : Improves performance with experience

- Has components
 - Learning element (improves with feedback)
 - Performance element (selects actions)
 - Critic (gives feedback)
 - Planner (suggests exploratory actions)



They modify their behaviour by observing the consequences of their actions, adjusting their internal models and decision-making approaches to achieve better outcomes in future interactions

Properties of Environments (PEAS)

P → performance measure

E → Environment

A → Actuators

S → Sensors

Components of Rational Agents

- Perception: The ability to perceive the environment through sensors
- Knowledge base: Information the agent has about the environment and itself
- Decision-making process: Algorithms and rules that guide the agent's actions
- Action: The ability to perform actions that affect the environment through actuators

→ Ex. for self-driving car

Design Principles of Rational Agents

- Performance measure: evaluates how well the agent is achieving its goals.
 - + Speed, safety, passenger comfort
- Rationality: acting optimally given the info and computational resources available.
 - slowing down when a pedestrian is crossing
 - Sensors → cameras, GPS, radars
 - Actuators → steering, brakes, accelerator
- Autonomy: can make decisions without human ~~intervention~~

Applications of Rational Agents

- Robotics (operate robots to perform tasks)
- financial trading (make buy-sell decisions)
- Health Care (assist in diagnosis & treatment planning)
- Game AI (controls non-player characters)
 - NPc

Early AI (Turing test) focused on mimicking humans.

- But AI today is built on rational agent approach.
- Instead of asking "Can machines think like humans?", it asks → "Can machines act rationally to achieve goals?"

Advantages over Turing Test

- | | |
|---|--|
| <ul style="list-style-type: none"> - Turing Test - focused on human-like conversation - subjective, imitation-based, limited | <ul style="list-style-type: none"> - Rational Agent Approach - General framework (not restricted to conversation) - Works for any environment (games, robotics, healthcare etc.) - Objective → maximizing measurable performance |
|---|--|

In AI, problem solving is often modeled as a search through a state space. Oh, that's ^{phewo.} what's we study we ^{are} studying searching algorithms.

State Space Representation

→ formal way of defining a problem as search

A State Space is the set of all possible states (configurations) that can be reached in the problem.
Problem solving → moving from the initial state to the goal state using operators (actions)

Components of a state space problem.

- Initial state → The starting point of the problem
- Goal state → Desired condition to be achieved
- State space → The set of all possible states reachable from initial state by applying actions.
- Actions → Rules for moving from one state to another.
- Path cost → Sequence of states from start to goal; path cost helps compare solutions

Ex. 8 puzzle problem

Initial state :
A scrambled
3x3 board
with tiles 1-8
and a blank
space

6	7	1
2	8	5
3		4

3x3

Goal state :
Tiles arranged
in order (1-8)
with blank space in bottom-right

1	2	3
4	5	6
7	8	

- Operators (Actions) : Move blank tiles UP/DOWN/LEFT/RIGHT
- State space : All possible arrangements of tiles ($9! = 362880$ states)
- Path cost : no. of moves

Advantages of State Space Representation

- Provides formal framework for problem solving
- Makes it easier to apply search algorithms (DFS, BFS, A*)
- Useful for heuristic design

State space representation defines a problem as search through states, from an initial state to a goal state, using operators, forming the basis of AI problem solving.

PAGE NO.
DATE

Limitations:

- State space may be too large (combinational explosion)
- Requires efficient search strategies (heuristic, pruning)

mit 2

UNINFORMED v/s INFORMED SEARCH

- Uninformed search (Blind search) (no heuristic)
- Search strategies that do not use any domain knowledge
- They only use the problem definition (initial state, action, goal state)

→ Explore the state space blindly until the goal is found

Ex BFS, DFS

- + Guaranteed to find optimal solution (if exist)
- + Simple to implement
- May explore huge parts of state space unnecessarily
- Very inefficient for large problems (high time, space complexity)
Time constan

- Informed Search (Heuristic Search)

- Search strategies that use problem-specific knowledge (heuristic) to guide search towards the goal
- Use a heuristic function $h(n)$ = estimated cost from node n to goal

Ex A* Search ($f(n) = g(n) + h(n)$), Best first search

- + Much more efficient (reduce search space) (less time, space complexity)
- + Often find optimal solution faster
- Requires a good heuristic (designing heuristics can be hard)
- Performance depends on accuracy of heuristic

State space representation defines a problem as search through states, from an initial state to a goal state, using operators, forming the basis of AI problem solving.

PAGE NO.
DATE

Limitations:

- State space may be too large (combinatorial explosion)
- Requires efficient search strategies (heuristic, pruning)

Unit 2

UNINFORMED v/s INFORMED SEARCH

- Uninformed search (Blind search) (no heuristic)
- Search strategies that do not use any domain knowledge
- They only use the problem definition (initial state, action, goal state)

→ Explore the state space blindly until the goal is found

Ex: BFS, DFS

- + Guaranteed to find optimal solution if exist)
- + Simple to implement
- May explore huge parts of state space unnecessarily
- Very inefficient for large problems. (high time, space complexity)
Time consuming

- Informed search (Heuristic Search)

Search strategies that use problem-specific knowledge (heuristic) to guide search towards the goal

→ Use a heuristic function $h(n)$ = estimated cost from node n to goal

Ex: A* Search ($f(n) = g(n) + h(n)$), Best first search

- + Much more efficient (reduce search space) (less time, space complexity)
- Often find optimal solution faster
- Requires a good heuristic (designing heuristics can be hard)
- Performance depends on accuracy of heuristic

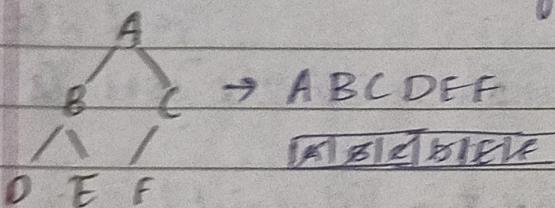
Unit 2

BFS (Breadth first search)

uninformed (blind) search algorithm which explores the state space level by level

Algorithm

1. Put the start node (root \rightarrow initial state) in queue (FIFO)
2. Repeat until queue is empty
 - \rightarrow remove from front node
 - \rightarrow if its goal state, return solution
 - \rightarrow otherwise, add all of its unvisited children to the queue



- Complete \rightarrow always find solution if exists
- Optimal \rightarrow give shortest path
- Time complexity: $O(b^d)$
- Space complexity: $O(b^d)$

$b \rightarrow$ branching factor
 $d \rightarrow$ depth of shallowest factor

(use case \rightarrow to find shortest path when state space is not too large (memory expensive))

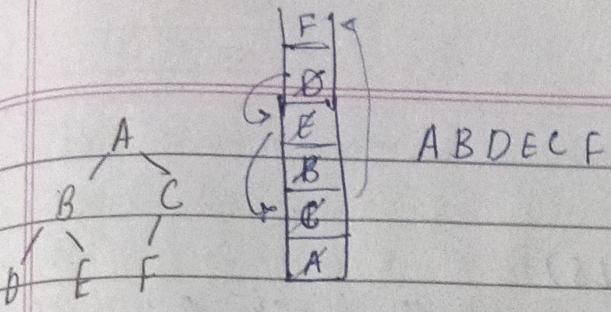
Unit 3

DFS (Depth first Search)

- uninformed (blind) search algorithm which explores a state space by going as deep as possible along one branch before backtracking

Algorithm

1. Put the start node (root \Rightarrow initial state) in stack (LIFO)
2. Expand the node:
 - \rightarrow if it's the goal state, stop
 - \rightarrow otherwise, visit the first unvisited child
3. If no children left, backtrack to explore other paths



- Incomplete (may get stuck in infinite paths in infinite depth trees)
- Not optimal (may find longer path before the shortest one)
- Time complexity : $O(b^d)$
- Space complexity : ~~$O(b^d)$~~ $O(bd)$

(use case → when memory is limited
when the search tree is very wide but not very deep)

unit 7

BIDIRECTIONAL SEARCH

- Runs two simultaneous searches :
 - One forward from initial state
 - One backward from goal state
- When two searches meet in the middle, a path is found
- Instead of searching depth d , each search goes about $d/2$, making it faster

A — B — C — D — E
start goal

forward search : A → B → C

Backward search : E → D → C

Meet at C, path found =

Algorithm

1. Initialize two frontiers (queues) → forward (start) & Backward (goal)
2. Expand nodes alternately from both frontiers
3. If a node is found in both frontiers → path is complete
4. Reconstruct solution by joining forward and backward paths

Optimal (if using BFS both sides)

Time complexity : $O(b^{d/2} + b^{d/2}) = O(b^{d/2})$

Space complexity : $O(b^{d/2})$

Disadvantage → need to know goal state in advance

BEST FIRST SEARCH

DEPTH LIMITED SEARCH (DLS)

DLS is a DFS with a cutoff depth limit l

- it explores nodes only up to a fixed depth l
- if the goal is beyond that depth, it will not be found
- ⊕ Helps avoid getting stuck in infinite paths (a problem in DFS).

Algorithm

1. Start at root node
 2. Perform DFS but stop if the current depth = limit l
 3. If goal is found within $\text{depth} \leq l$, return solution
 4. Otherwise return cutoff (meaning deeper search required)
 - ~~is~~ not complete, if limit < depth of solution
 - Not optimal
 - Time Comp.: $O(b^l)$
 - Space Comp.: $O(bl)$
- + prevents infinite loop (unlike DFS)
- + less memory than BFS
- If limit is too small → may miss the sol.
- If limit is too large → behave like DFS

Heuristic Search Approach

A Heuristic is a problem searching specific rule or function that provides an estimate of how close a given state is to the goal

→ $f(n) = \text{estimated cost from node } n \text{ to goal}$
It does not guarantee accuracy, but helps guide the search efficiently.

Heuristic search is a search strategy that uses heuristic information (domain knowledge) to prioritize which path to explore

Misplaced tiles
Manhattan dist
Euclidean dist
Chebyshev dist

Unlike uninformed search (blind), heuristic search is goal-directed.

Heuristic Algorithms

Greedy Best First Search

$f(n) = h(n) \rightarrow$ Manhattan dist.)

A* Search $f(n) = g(n) + h(n)$)

Beam Search, Hill Climbing

Properties

- Complete (A* is complete, hill climbing is not)
- Optimal (depends on heuristic)
- Time complexity (less than uninformed search, but depends on heuristic quality)
- Space complexity (often high \rightarrow A* stores all nodes in memory)

Advantages

- Much more efficient than uninformed methods
- finds solution faster by focusing on promising paths

Disadvantages

- Requires good heuristic function (designing one can be hard)
- Bad heuristic can make search worse than uninformed search.

ITERATIVE DEEPENING SEARCH (Iterative deepening DFS)

IDS is a hybrid of BFS and DFS.

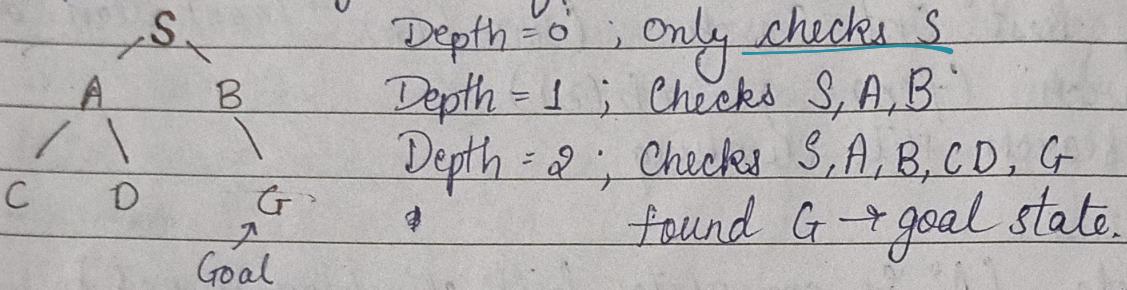
- it performs depth-limited search (DLS) repeatedly with increasing depth limits

- at depth limit 0, 1, 2... d until the goal is found.

→ It combines the space efficiency of DFS with the completeness & optimality of BFS

Working

1. Start with depth limit = 0 \rightarrow run DLS
2. If goal not found, increase depth limit to 1 \rightarrow run DLS again
3. Repeat until goal is found.



Properties

Complete (like BFS)

Optimal {

Time complexity : $O(b^d)$, same as BFS

Space complexity : $O(bd)$, same as DFS (much better than BFS)

Advantages

- Uses much less memory than BFS
- Guarantees finding the optimal solution (unlike DFS)
- Works well in large, infinite, or unknown depth problems.

Disadvantages

- Some nodes are re-expanded multiple times, which increase overhead.
- But since most nodes are near the deepest level, extra cost is not too high

BEST FIRST SEARCH (heuristic search)

It is a heuristic search that explores the most promising node first, based on an evaluation function:

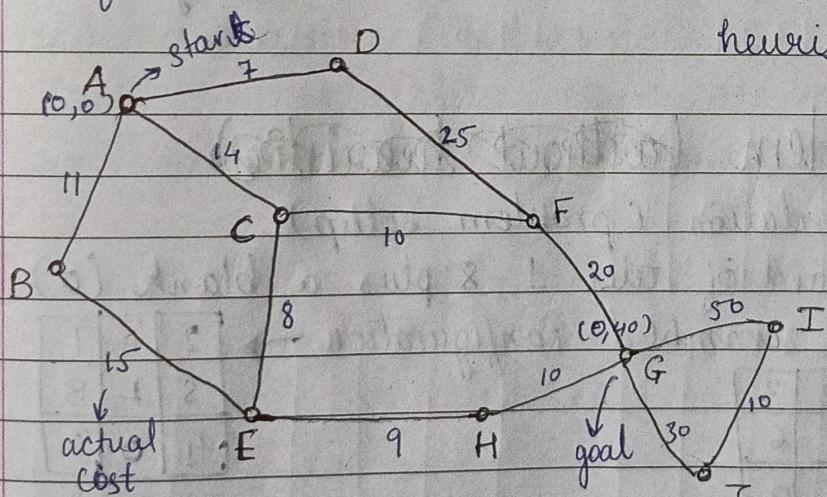
$$f(n) = h(n)$$

where, $h(n)$ = heuristic estimate of cost from node n to goal
 → It expands the node that appears to be closest to the goal

Algorithm

1. Put the start node A into a priority queue
2. While the queue is not empty.

- Remove the node with lowest heuristic value $h(n)$
- If its goal \rightarrow stop.
- Else expand its children and add them to the priority queue.



heuristic func = manhattan dist

$$h(x) = \sqrt{(x_1 - x_2)^2 + (y_1 - y_2)^2}$$

$$A \rightarrow G = 40$$

$$B \rightarrow G = 32$$

$$C \rightarrow G = 25$$

$$D \rightarrow G = 35$$

$$E \rightarrow G = 19$$

$$F \rightarrow G = 17$$

$$G \rightarrow G = \text{closed} \quad (Goal)$$

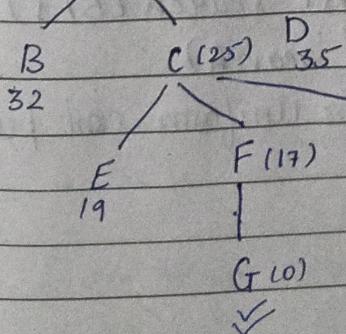
$$H \rightarrow G = 10$$

priority = A C F G
queue

$$A(40) \rightarrow h(x)$$

path = A → C → F → G

$$\text{cost} = 14 + 10 + 20 = 44$$



other possible paths (without heuristic)

$$A \rightarrow B \rightarrow E \rightarrow H \rightarrow G$$

$$\text{cost} = 41$$

- Complete (if graph is finite)
 - Optimality not guaranteed (may choose a locally good path, but miss global optimum)
 - Time Complexity : $O(b^d)$ worst case
 - Space Complexity : Keeps all nodes in memory (can be large)
- + faster than blind searches (BFS/DFS)
- + Uses heuristic, so guided towards goal
- Not always optimal
- Depends heavily on the quality of heuristic

Best first Search = Greedy heuristic search → expands node with the lowest $h(n)$ first
 "always go where it looks closest to the goal."

8-puzzle problem ~~(without heuristic)~~

→ State Space representation (problem setup).

State: a 3×3 board of tiles 1-8 plus a blank (0)

Initial state: any scrambled configuration →

1	2	3
4	5	6
7	8	0

2	6	1
5	1	3
4	0	8

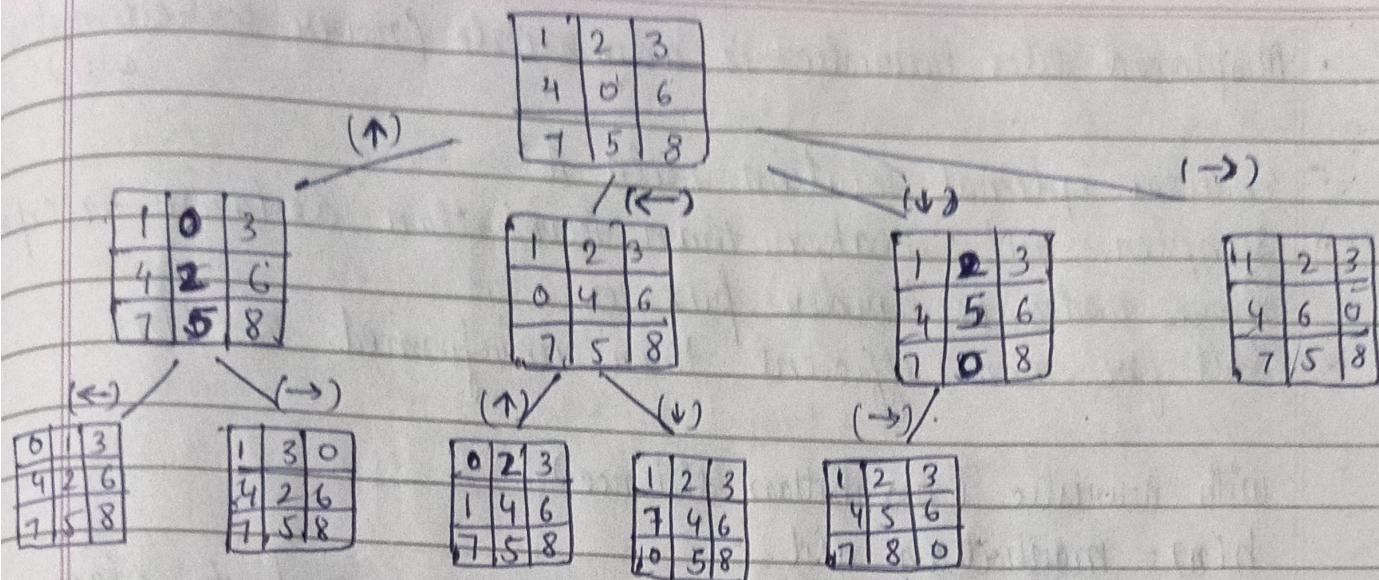
Actions: move blank UP(↑), DOWN(↓), LEFT(←), RIGHT(→)

Path cost: 1 per move (shortest path are optimal)

without heuristic → use Breadth first search (BFS)
 (uninformed search)

BFS guarantees shortest solution (for uniform cost per step)

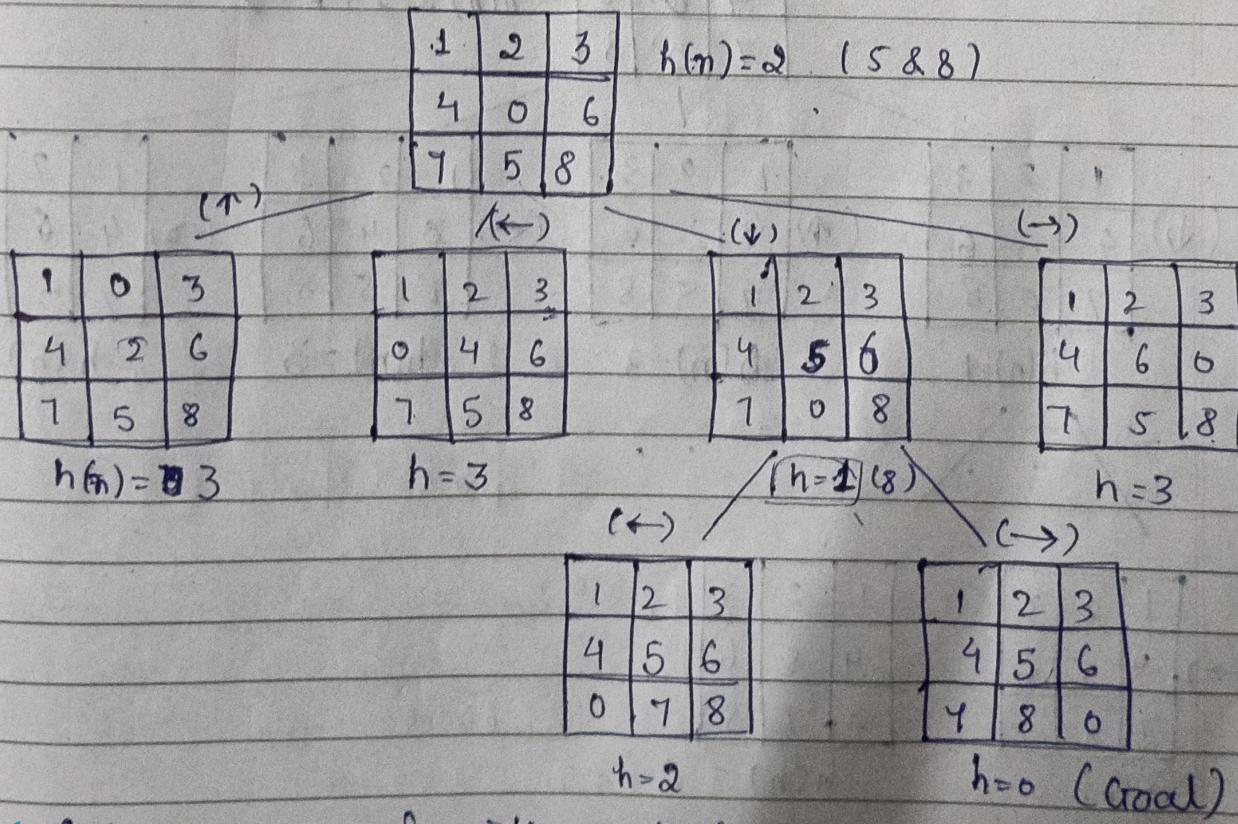
BFS may consume a lot of memory



If the initial state is far from goal (large depth), BFS may be infeasible due to memory and time. Thus we use heuristic (A*)

with heuristic (Misplaced Tiles)

Count the no. of tiles which are not in correct place, and will expand the branch with lowest no. of misplaced tiles



faster than non heuristic method

- ✓ Misplaced tiles heuristic is admissible (never overestimates cost)
- ✓ Ensures optimal solution with A*
- ✓ Simpler but weaker than Manhattan distance. → expands more nodes in harder puzzles.
- ✓ Still far more efficient than uninformed BFS/DFS

with heuristic (Manhattan distance)

$h(n) = \text{manhattan dist}$

↳ sum of the moves required to place a misplaced tile in right position.

2	1	3
5	6	4
7	8	0

2 → 1 move

1 → 1 move

3 → 0 move

8 → 1 move

6 → 1 move

4 → 2 move

7 → 0 move

8 → 0 " "

$$h(n) = 6$$

$$h(n) = 1 + 1 + 0 + 1 + 1 + 2 + 0 + 0 + 0 = 6$$

1	2	3
4	0	6
7	5	8

$$h(n) = 1 + 1 = 2$$

1	2	3
4	5	6
7	0	8

$$h(n) = 1$$

1	0	3
4	2	6
7	5	8

$$h(n) = 3$$

1	2	3
0	4	6
7	5	8

$$h(n) = 3$$

1	2	3
4	5	0
7	5	8

$$h(n) = 3$$

1	2	3
4	5	6
7	8	0

$$h(n) = 0$$

1	2	3
4	5	6
0	7	8

$$h(n) = 2$$

Goal

manhattan dist: counts total number of horizontal + vertical steps each tiles is away from its goal.

PAGE NO.

DATE

- Manhattan is generally stronger, so it typically expands fewer nodes, making the search faster on harder problems
- Complexity: exponential (worst-case)
- Always admissible (never overestimates) \rightarrow ensures optimal solution

$$h(n) = \sum_{i=1}^g |r_{\text{current}}(i) - r_{\text{goal}}(i)| + |c_{\text{current}}(i) - c_{\text{goal}}(i)|$$

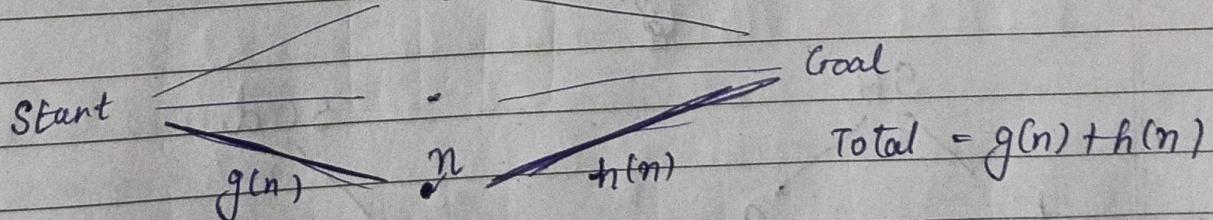
A* Search (informed searching) (heuristic search)

A* is a best-first search algorithm that finds the shortest path to a goal.

It combines:

- Uniform Cost Search: explores cheapest path so far $\rightarrow g(n)$
- Greedy Best-first search: explores the node closest to the goal $\rightarrow h(n)$

$$f(n) = g(n) + h(n)$$



$g(n)$: cost from start to node n

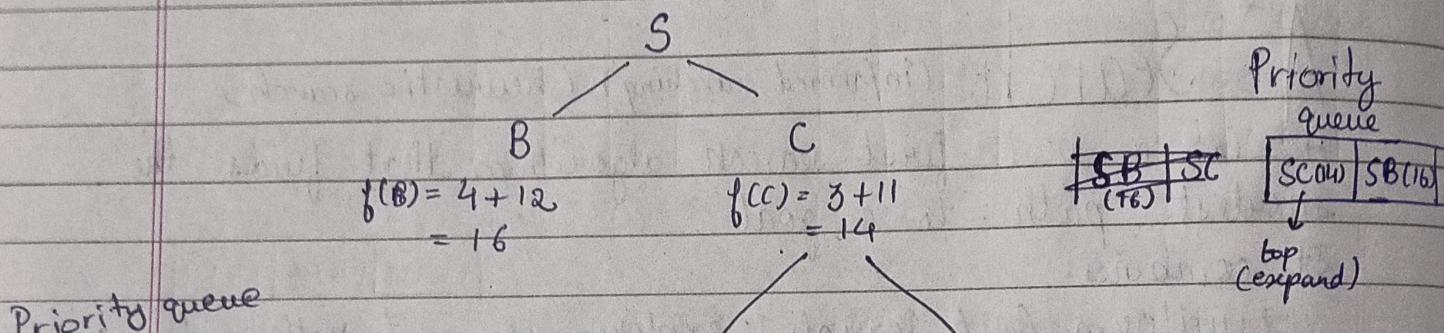
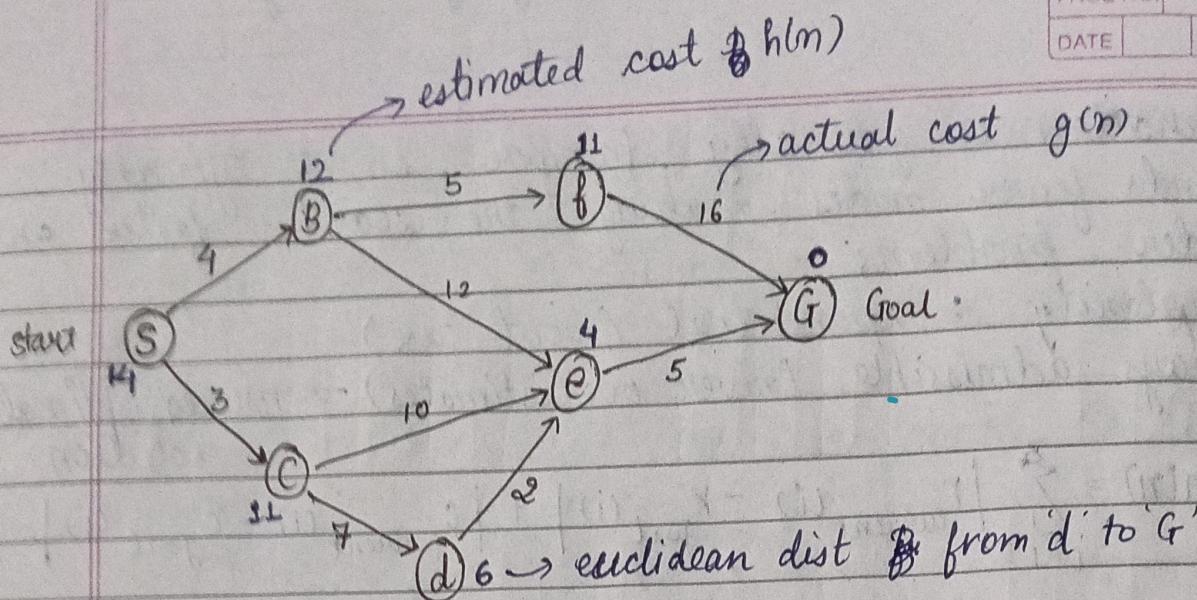
$h(n)$: estimated cost from n to goal (heuristic)

$f(n)$: estimated total cost of solution through n .

- Optimal (if $h(n)$ is admissible, never overestimates)
- Complete (if branching factor is infinite)
- Efficient when heuristic is strong.

F1 F2 F3 admissible \rightarrow never overestimates.

PAGE NO.	
DATE	



SB

$$f(f) = (9) + 11 = 20$$

$$f(e) = 16 + 4 = 20$$

scd

$$e \quad f(e) = (3 + 7 + 2) + 4 = 12 + 4 = 16$$

PQ = [Scde(16) | Scce(17) | SBf(20) | SBe(20)]

\downarrow

expand

Scde

1

G

$$f(n) = (3 + 7 + 2 + 5) + 0 = 17$$

path: $S \rightarrow C \rightarrow D \rightarrow E \rightarrow G$ (cost = 17)

\downarrow

skip the one, because $S \rightarrow e$ is already discovered with lower cost

A* Algorithm

1. Put the start node in OPEN list (priority queue sorted by f)
2. Remove the node with smallest $f(n)$ from OPEN
3. If it's goal \rightarrow stop, return path
4. Otherwise, expand it :
 - \rightarrow Generate successors (valid moves)
 - \rightarrow Compute $f = g + h$ for each
 - \rightarrow Add them to OPEN (skip if already discovered with lower cost)
5. Repeat until goal found or OPEN is empty

Advantages

- ✓ Guarantees optimal solution (with admissible h)
- ✓ Explores fewer nodes than BFS / UCS if heuristic is strong

Limitations

- ✓ Can use a lot of memory (keep all nodes in OPEN)
- ✓ Exponential in worst-case
- ✓ Heuristic quality is crucial (poor $h \rightarrow$ degenerates to UCS).

8 puzzle with A*

$$f(n) = g(n) + h(n)$$

\hookrightarrow manhattan dist

Initial state = $(1, 2, 3, 4, 5, 6, 0, 7, 8) \rightarrow$

1	2	3
4	5	6
0	7	8

Goal state = $(1, 2, 3, 4, 5, 6, 7, 8, 0)$

Actions = UP, DOWN, RIGHT, LEFT

State space = 9!

Path cost = 1 per move.

$$\text{heuristic} = \text{manhattan dist} = \sum_{i=1}^8 |r_{\text{curr}}(i) - r_{\text{goal}}(i)| + |c_{\text{curr}}(i) - c_{\text{goal}}(i)|$$

open = $A(2)$

(A)

1	2	3
4	5	6
0	7	8

$$f = 0 + 2 = 2$$

PAGE NO.

DATE

$g = 1$

(B)

1	2	3
0	5	6
4	7	8

$$f = 1 + 3 = 4$$

(C)

1	2	3
4	5	6
7	0	8

$$f = 1 + 1 = 2$$

open = $G(2)$ $B(4)$

$g = 2$

(D)

1	2	3
4	0	6
7	5	8

$$f = 2 + 2 = 4$$

open = $E(2)$ $B(4)$ $D(4)$

pop it

→ Goal reached → Stop

(E)

1	2	3
4	5	6
7	8	0

$$\begin{aligned} f &= 2 + 0 \\ &= 2 \end{aligned}$$

(GOAL)

$$f = 2 + 2 = 4$$

SKIP to push in
OPEN

This state is already
discovered with
lower cost

X

BEAM SEARCH (heuristic search)

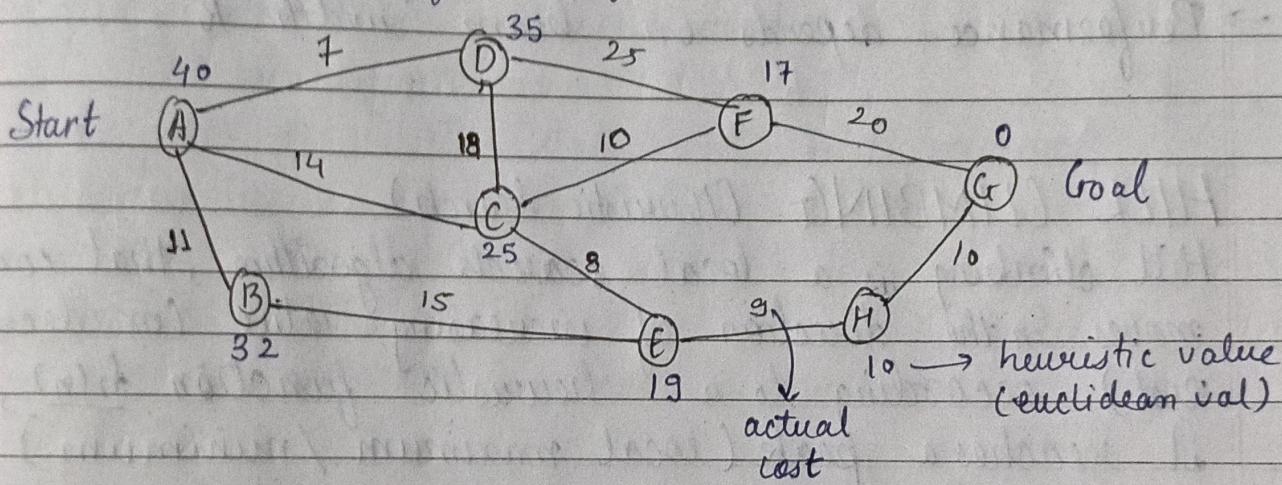
It is a heuristic search algo. that is a memory efficient version of Best-first Search

- At each level of the search tree, instead of keeping all generated nodes, it keeps only the best k -nodes (based on heuristic value)
- Here, k = beam width (a fixed integer)

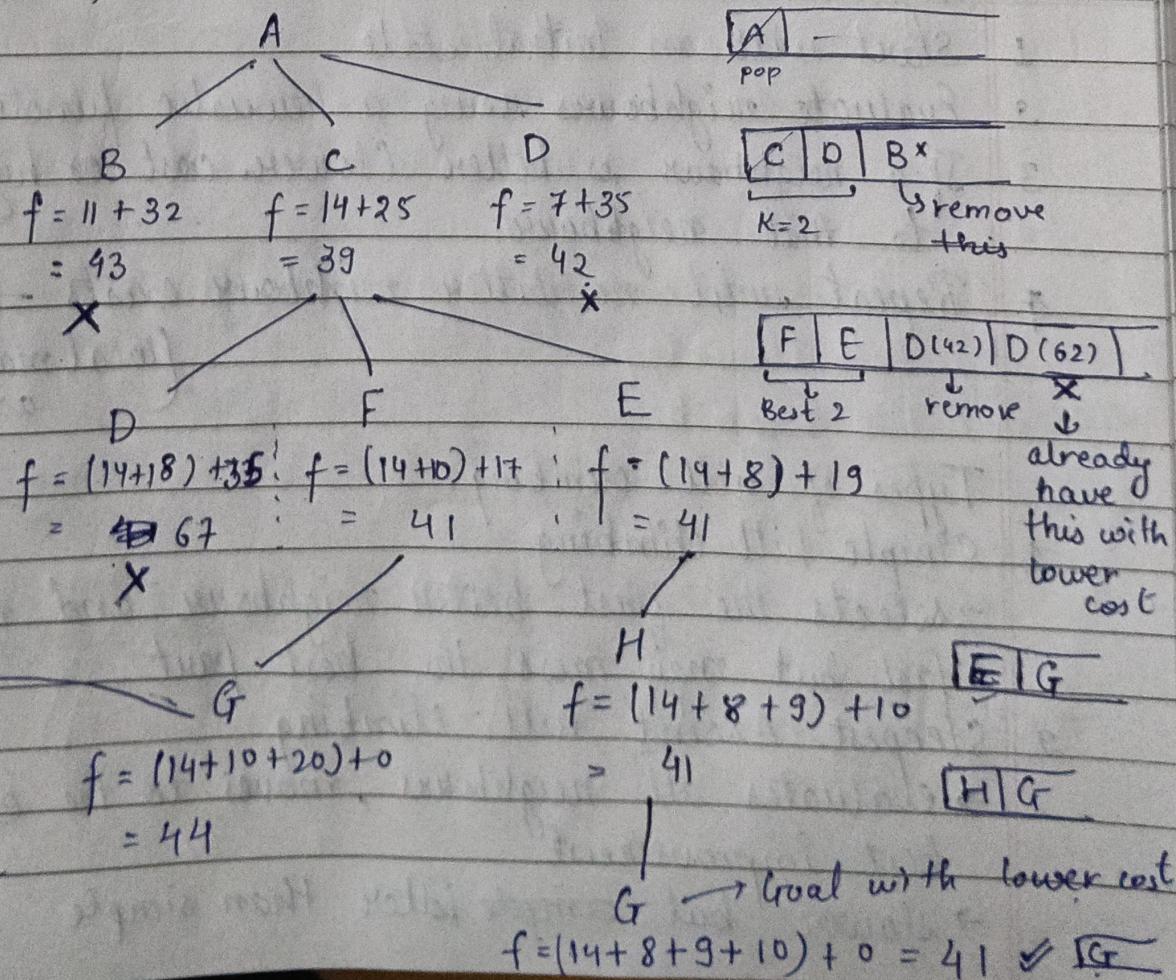
This reduces memory usage and search time but risks losing the optimal solution if the correct path is discarded

Algorithm

1. Start from initial node
2. Generate successors (children)
3. Evaluate them using a heuristic function $h(n)$
4. Keep only the k best nodes (lowest h values) and discard others.
5. Expand those nodes in the next level
6. Repeat until the goal is found (or no nodes remain)



Beam width, $k = 2$



- ✓ : Not complete (it may discard the path that leads to goal)
 - ✓ : Not optimal (may miss the best path)
 - ✓ : Space complexity : $O(K \cdot d) \approx \text{constant}$
 - ✓ : Time complexity : $O(K \cdot d \cdot d)$
- $(d = \text{depth of tree})$
 $(b = \text{branching factor})$
- ✓ → Uses less memory than Best first or A*
 - ✓ → faster because few nodes are expanded
 - ✓ → Works well when exact optimality is not critical
 - ✓ → Performance depends on beam width k.

HILL CLIMBING (Heuristic Search)

Hill climbing is a local search algorithm, that continues moves in the direction of increasing value (or decreasing cost) according to a heuristic function $h(n)$, until it reaches a peak (local maximum / minimum)

Algorithm

1. Start with an initial state
2. Evaluate neighbours using a heuristic function
3. If a neighbour is better (lower cost or higher value), move to that neighbour
4. Repeat until no better neighbour exists → stop
(local maximum / minimum achieved)

Types of hill climbing

1. Simple hill climbing
 - selects the first better neighbour and moves there.
 - fast but may miss the best part.
2. Steepest-Ascent hill climbing
 - evaluates all neighbours, moves to the one with the best improvement.
 - slower but ~~simpler~~ better than simple.

Hill climbing = Beam search ($k=1$)

PAGE NO.		
DATE		

Stochastic hill climbing

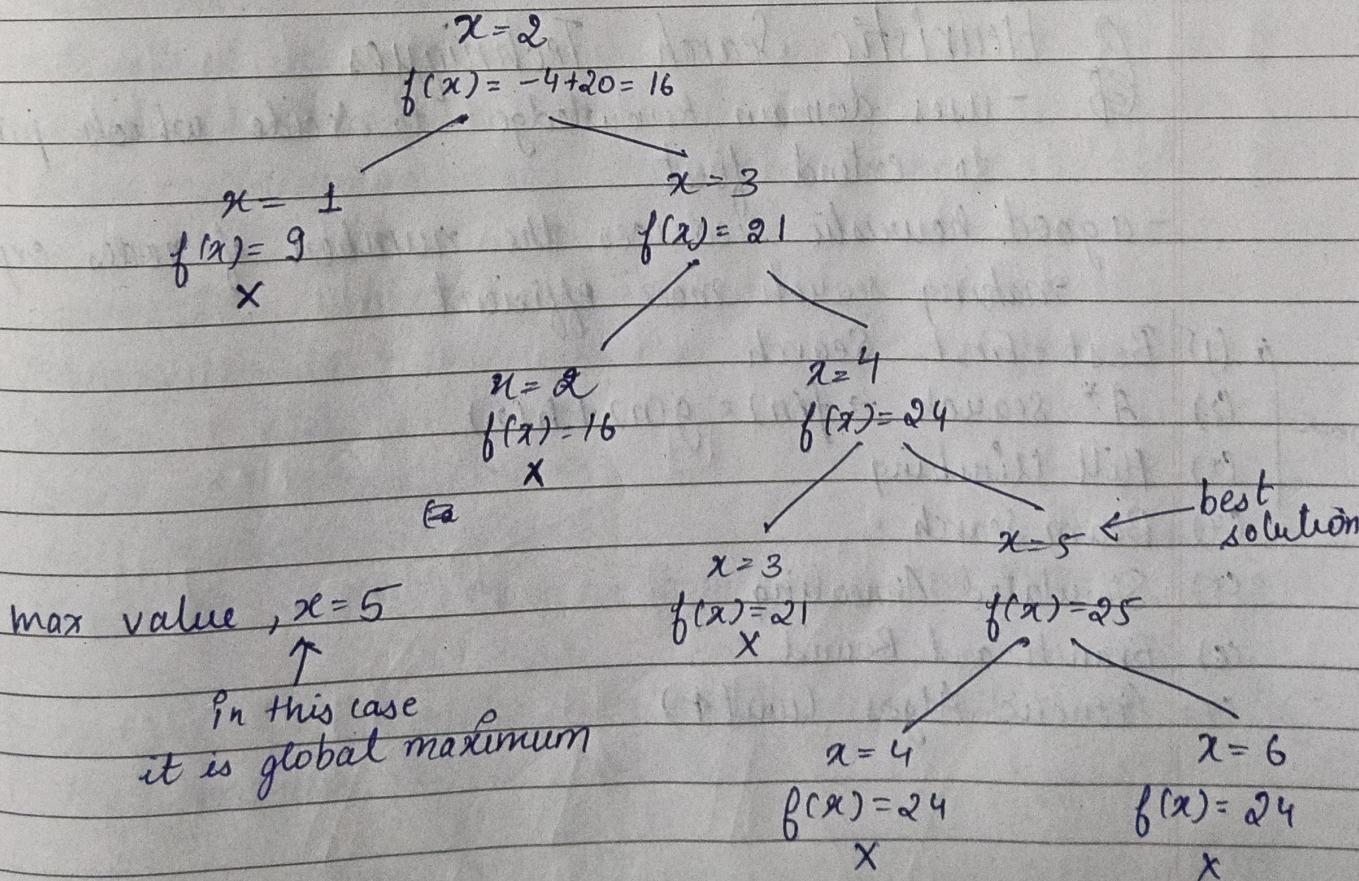
- chooses randomly among better neighbors (not always the best one)
- helps escape loops

Problems with hill climbing

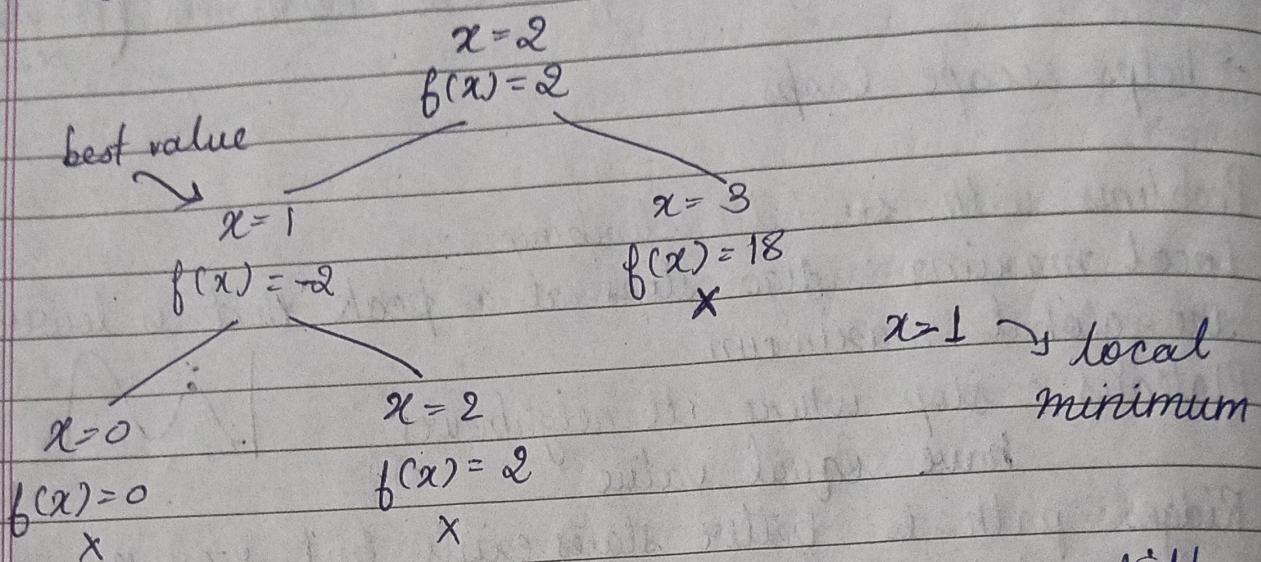
- Local maxima : Algo stops at a peak that is lower than the global maximum
- Plateaus : stop where all neighbours have equal value
- Ridges : path to better states exists but requires moving away temporarily.

Simple hill climbing

$$f(x) = -x^2 + 10x \rightarrow \text{maximize value.}$$



$$f(x) = x^3 - 3x \rightarrow \text{minimize}$$



If the global optimum were further away; hill climbing would never reach it
 It simply halts at first peak or valley \rightarrow (equal vals)

Heuristic Search Techniques

- uses domain knowledge to decide which path to extend first
- a good heuristic reduces the number of nodes explored making search more efficient.

- 1) Best-First Search
- 2) A^* Search $\rightarrow f(n) = g(n) + h(n)$
- 3) Hill Climbing
- 4) Beam Search
- 5) Simulated Annealing
- 6) Branch and Bound
- 7) Genetic Algos (Unit 4)

have adversarial search problems.

PAGE NO.	
DATE	

GAME PLAYING

- Game playing in AI is the process of building systems that can play strategic games against a human or another AI.
- These are multi-agent environments
 - Player 1 (maximizer)
 - Player 2 (Minimizer)

Types of Games

1. Single player games

- Only one agent, no opponent

Ex: 8-puzzle, Sudoku

Algos → search Algos (BFS, DFS, A* etc.)

2. Two-player games

- Involved competition

Ex: Chess, Checkers, Tic-Tac-Toe *

Algos → Minimax algo, Alpha-Beta pruning

Game tree / search tree / space search graph

- Representation of all possible moves in a game.
- Nodes → states of game
- Edges → actions / moves
- ply → depth of graph
- Terminal nodes → win / loss / draw state

→ We use a heuristic evaluation function to estimate the "goodness" of a position

Ex: Deep Blue (Chess)

Alpha Go (Go)

Minimax evaluates all possible moves and outcomes

- Max (X) chooses best utility move
- Min (O) blocks / reduce score

Backtracking → This ensures optimal play

PAGE NO.	
DATE	

Minimax Algorithm (adversarial search)

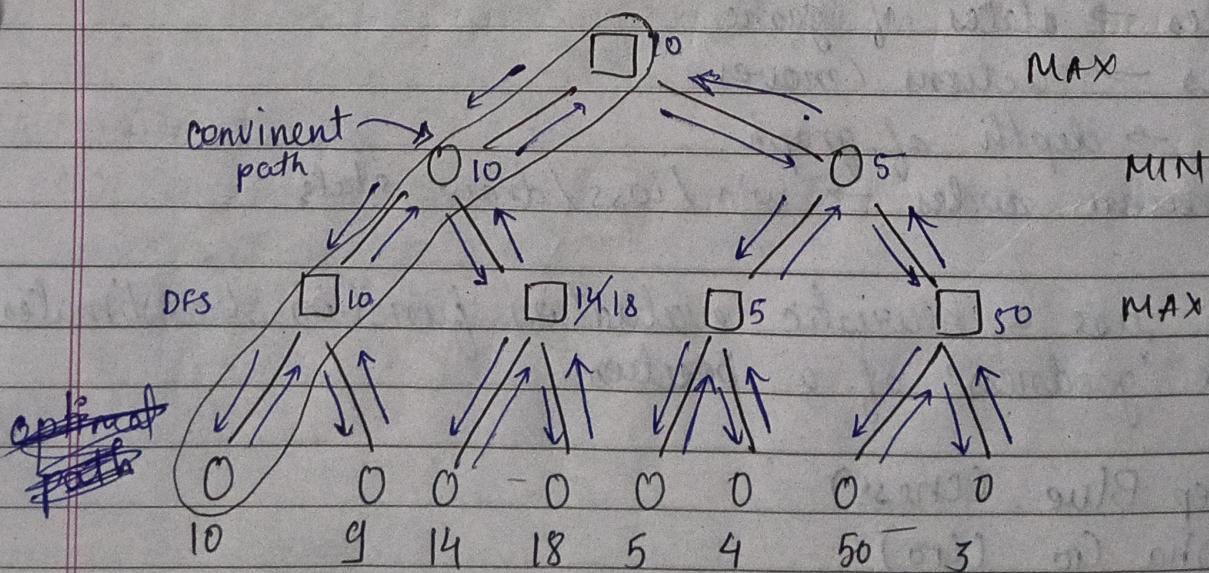
- Decision-making algo used in 2-player games
- It assumes
 - Player 1 (Maximizer) tries to maximize the score
 - Player 2 (Minimizer) tries to minimize the score
- Both players are rational and play optimally

Working

- Build a game tree of all possible moves
- Assign utility values (+1 → win, -1 loss, 0 draw)
- Backtrack values up the tree
 - At MAX nodes → choose the maximum value
 - At MIN nodes → choose the minimum value

Algorithm

1. Generate the game tree to a certain depth (or terminal states)
2. Apply an evaluation function at terminal nodes.
3. Backtrack values
 - MAX node = choose maximum child
 - MIN node = choose minimum child
4. Root node's value = best move for the current player



MAX aims to maximize its score, while MIN aims to minimize it.

PAGE NO.	
DATE	

Limitations

Exponential complexity $O(b^d)$

$b \rightarrow$ branching factor

$d \rightarrow$ depth

Not feasible for large games like chess \rightarrow use Alpha-Beta

MINIMAX FOR TIC-TAC-TOE

Player X (AI) \Rightarrow MAX

Player O ~~Human~~ \Rightarrow MIN

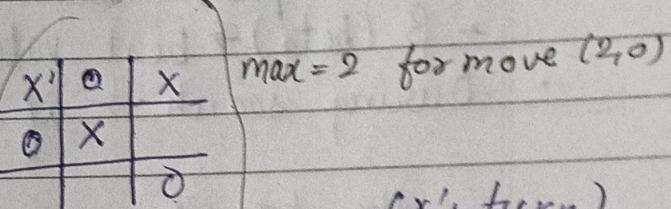
X wins $\rightarrow +10$

O wins $\rightarrow -10$

Draw $\rightarrow 0$

Utility values

1. Start from current board



max = 2 for move (2,0)

2. Generate all possible moves

3. Backtrack

max = 0

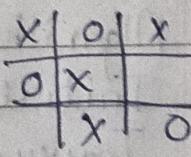
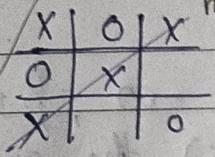
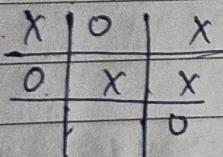
(2,2)

(2,0)

(X's turn)

(2,1)

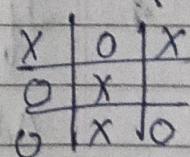
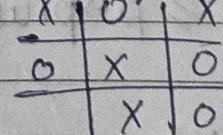
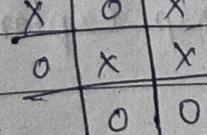
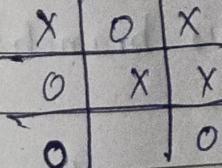
(MAX)



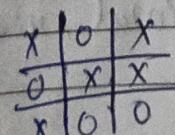
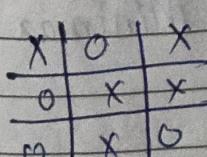
win for X
utility = +10

O's turn

(MIN)



X's turn
MAX



Draw
utility = 0

Draw
utility = 0

Draw
utility = 0

Draw
utility = 0

optimized version of minimax

Alpha-Beta Pruning

- Optimization of minimax algo
- It eliminates branches in the game tree that cannot affect the final decision
- This makes the search faster without changing the result.

Keep track of two values while searching

$\rightarrow \alpha$ (alpha) \rightarrow best value that MAX can guarantee so far.

$\rightarrow \beta$ (beta) \rightarrow best value that MIN can guarantee so far.

- During search \rightarrow if at any point;

$$\alpha > \beta$$

the branch is pruned (cut off) because it won't influence the final decision

Algorithm

1. Initialize $\alpha = -\infty$, $\beta = +\infty$

2. Traverse (DFS) the game tree like minimax

3. At each node

- update α at MAX nodes

- update β at MIN nodes

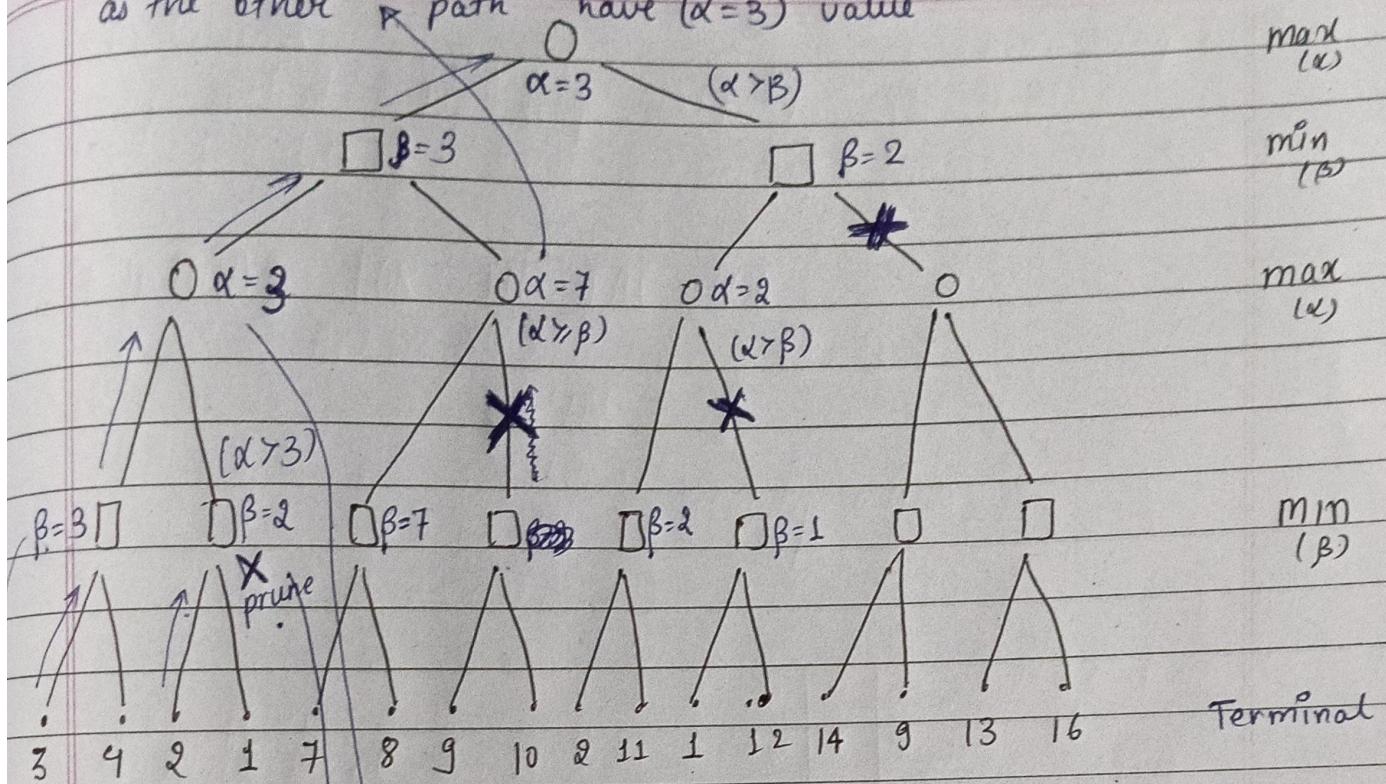
4. If $\alpha \geq \beta \rightarrow$ prune (skip exploring this branch)

Advantages

- Reduce node explored from $O(b^d)$ $\rightarrow O(b^{d/2})$ in best case.
- Efficient in deep game tree.
- Guarantee same optimal result as Minimax

no need to explore this branch further
as the MIN node will not choose this path
as the other path have ($\alpha = 3$) value

PAGE NO.		
DATE		



~~MIN~~ value
of this node
will not exceed
 3 (at max 3)

the MAX value of
this node will not
be less than 3
(at least 3)

no need to ~~explore~~
explore this branch
further because it
can give at ~~most~~ max
value of 2 (less than current α).

